

HARSH NANGLIA

+91 8950123796 | Nangliaharsh@gmail.com | linkedin.com/in/harsh-nanglia

PROFESSIONAL SUMMARY

AI/ML Engineer and Python Developer with hands-on experience building production GenAI systems, agentic workflows, and document intelligence pipelines on Azure. Shipped solutions across EMS dispatch, multi-modal document extraction, and RAG-based retrieval — with measurable results: 70% reduction in manual data entry, 50% drop in API latency. Works end-to-end across LLM orchestration, backend engineering, and cloud deployment.

TECHNICAL SKILLS

Generative AI & LLMs: LangChain, LangGraph, RAG (hybrid search, HyDE, re-ranking), Agentic AI, Prompt Engineering, Claude API, Fine-tuning, OpenAI / Azure OpenAI, Google Gemini, Hugging Face Transformers, Vector Databases

ML / Deep Learning: TensorFlow, PyTorch, Scikit-learn, NLP, Model Evaluation (RAGAS, LLM), Feature Engineering

Cloud & MLOps: Microsoft Azure (Functions, Key Vault, OpenAI, Cognitive Services, Document Intelligence, Maps), Firebase, Docker, REST APIs, CI/CD

Backend, CV & OCR: Python, FastAPI, Flask, Streamlit, pyodbc, JavaScript, Azure Vision API, Tesseract, OpenCV, pdfplumber, Camelot, PyPDF

Data & Databases: SQL Server, PostgreSQL, MongoDB, MySQL, Pandas, NumPy, Power BI

PROFESSIONAL EXPERIENCE

Python Developer – AI & Automation

Apr 2025 – Oct 2025

Digital Glyde | Full-time, Onsite

ETHAN – Emergency TeleHealth & Navigation System (U.S. EMS)

- Designed and built the **FastAPI + Azure** backend for real-time EMS dispatch — owning API contracts, database schemas, and deployment end-to-end; wrote a regression suite that caught 15+ bugs before release.
- Cut **API latency by 50%** by re-indexing hot SQL Server tables and rewriting N+1 query patterns into parameterised pyodbc batches.
- Integrated **Azure Maps** for live ambulance tracking, ETA calculation, and route optimisation across concurrent multi-incident scenarios.

IDE – Intelligent Document Extraction

- Built a **multi-modal extraction pipeline** across 100+ heterogeneous PDFs and Excels, replacing manual data entry and reducing per-document processing time by ~70%.
- Composed **Azure OpenAI, Document Intelligence, Tesseract, and OpenCV** under a strict JSON schema; validated spatial correctness with Plotly bounding-box overlays.
- Improved extraction accuracy on noisy scanned documents using targeted prompt engineering and few-shot grounding.

Personal R&D | Independent AI Engineering & Research

Nov 2025 – Present

DocuMind | RAG-Powered Document Intelligence System

- Built a production-grade RAG pipeline using **LLaMA 3, ChromaDB, Langchain and FastAPI** enabling natural language **Q&A** across uploaded documents with cited responses. Measured retrieval accuracy via RAGAS evaluation framework.
- Optimized retrieval quality through chunking strategy tuning, overlap experimentation, and system prompt iterations, reducing hallucination to near-zero on domain-specific documents..
- Deepening expertise in **Agentic AI**, advanced RAG architectures, and **LLM fine-tuning** through structured self-study and hands-on experimentation.

DATAPLAY, Jaipur (Remote) | Data Analyst Intern

June 2024 – July 2024

Designed interactive dashboards in **Power BI**, increasing decision accuracy by 35%. **Led a team of four** to deliver a **Data visualization** project under tight deadlines. Conducted data analysis to identify trends, driving cost optimization.

PROJECTS

KAIROS AI | Multi-Tool Generative AI Platform

Mar 2024 – Nov 2024

- Built a unified **GenAI** platform combining **conversational AI**, multi-PDF Q&A, **voice assistance**, and language utilities — served via a JavaScript + REST API frontend.
- Implemented a **chat-with-multiple-PDFs** module backed by **OpenAI** embeddings and a retrieval layer, enabling natural-language exploration of large document corpora with follow-up queries.

AI News Video Generation Pipeline | LLM-Powered Content Automation

May 2025 – Jun 2025

- Designed an end-to-end pipeline that converts trending topics into short-form video assets via **Google Gemini** (script) and **Imagen** (visuals); uses structured prompt templates and **NLP** post-processing for tone/shot-list consistency, with **Firebase Firestore** for session persistence.

EDUCATION

B.Tech, Computer Science and Engineering

Jun 2021 – Jun 2025

Maharaja Agrasen Institute of Technology (MAIT), Baddi | CGPA: 8.50